CS 6210/CS 4210 ADVANCED OPERATING SYSTEMS	Fall 2017
Distributed Shared Memory (Part I)	
Prof. Umakishore Ramachandran	Scribe: Long Gong

1 Overview

In this lecture we cover another way to exploit remote memories, namely software implementation of distributed shared memory (DSM). That is to create an operating system abstraction that provides an illusion of shared memory to the applications, even though the nodes in the local area network do not physically share memory.

DSM asks the question, *if shared memory makes life simple for application development in a multiprocessor system, can we try to provide the same abstraction in a distributed system, and make the cluster look like a shared memory machine ?*
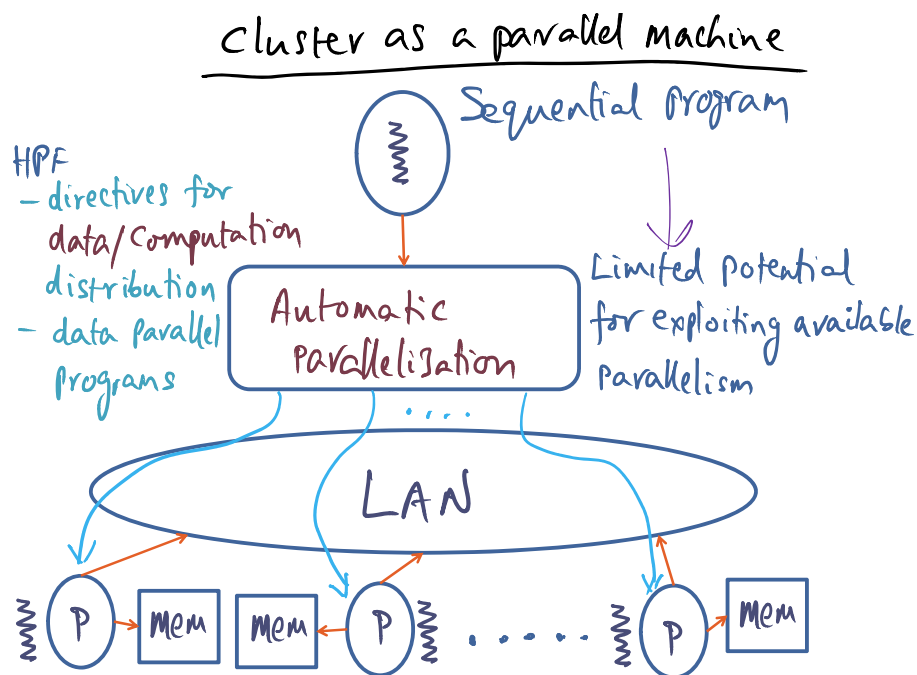
2 Motivation: Cluster as a Parallel Machine

Figure 1: Automatic parallelization.

Automatic Parallelization. Assume that if the starting point is a sequential program. One possibility is to use so-called **automatic parallelization**, an illustration of which is shown in Figure 1. That is, instead of writing an explicitly parallel program, we write a sequential program. And let somebody else do

---

[1]The first several sections of this scribe are actually covered by a lecture before this one. They are included for this scribe to be self-contain.

[2]All the figures in this scribe are directly from or modified from the lecture of Prof. Umakishore Ramachandran.

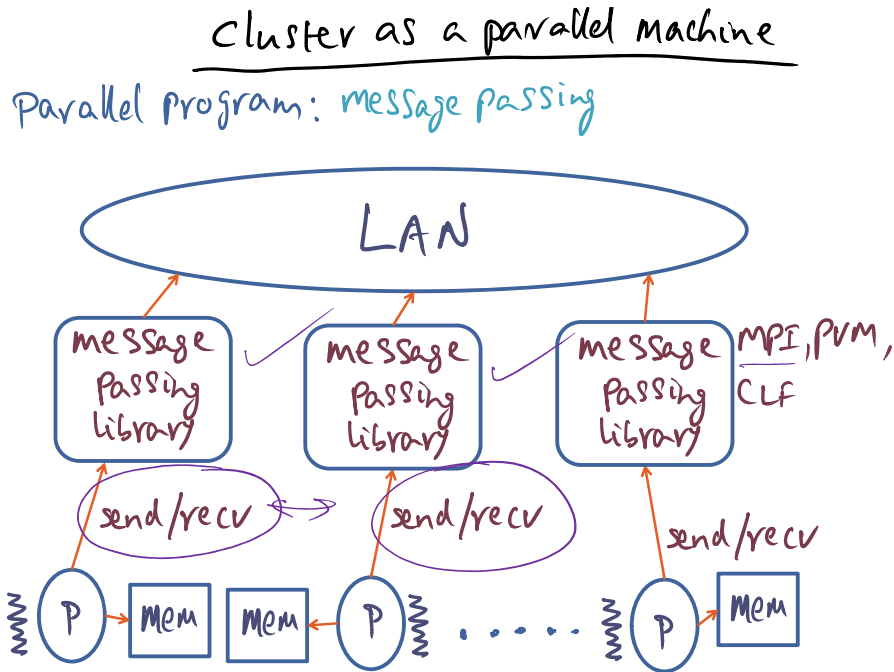## Cluster as a parallel machine

Parallel program: message passing

Figure 2: Message passing style of parallel programming.

the heavy lifting in terms of identifying opportunities for parallelism that exits in the program and map it to the underlying cluster. That is called an implicitly parallel program. There are opportunities for parallelism, but the program itself is not written as a parallel program. Now it is the onus of the tool, in this case, an automatic parallelizing compiler, to look at the sequential program and identify opportunities for parallelism and exploit that by using the resources that are available in the cluster. High-performance FORTRAN is an example of a programming language that does automatic parallelization, but it is user-assisted parallelization in the sense that the user who is writing the sequential program is using directives for distribution of data and computation. And those directives are then used by this parallelizing compiler for mapping these computations onto the resources of a cluster. So it puts it on different nodes on the cluster and that way, it exploits the parallelism that is there in the hardware, starting from the sequential and doing the heavy lifting in terms of converting the sequential program to a parallel program to extract performance for this application. This kind of automatic parallelization, or implicitly parallel programming, works really well for certain classes of program called data parallel programs. In such programs, for this most part, the data accesses are fairly static, and it is determinable at the compile time. So in other words, there is limited potential for exploiting the available parallelism in the cluster if we resort to implicitly parallel programming.

Message Passing. Therefore, we write the program as a truly parallel program instead of a implicitly parallel program. In other words, the application programmer is going to think about his/her application and write the program as an explicitly parallel program. And there are two styles of writing explicitly programs. And correspondingly system support for those two styles of explicitly parallel programs. One is called **message passing** style of explicitly parallel program. The run time system is going to provide a message passing library which has primitives for an application thread to do sends and receives to its peers that are executing on other nodes of the cluster. Figure 2 gives an illustration of message passing style of explicitly parallel program. So this message passing style of explicitly parallel program is true to the physical nature of the cluster. The physical nature of the cluster is the fact that *every processor have its private memory and this memory is not shared across all the processors.* So the only way a processor can communicate with another processor is by sending a message through the network from which one processor can receive messages. The processor can not directly reach into the memory of the other processor because that is not the way the cluster is architected. So the message passing library is true to the physical nature

Figure 3: DSM style of parallel programming.

of the cluster that there is no physically shared memory. And there are lots of examples of message passing libraries that have been written to support explicitly parallel programming in a cluster, such as message passing interface (MPI), Parallel Virtual Machine (PVM) [1], CLF [2]. And to this day, many scientific applications running on large-scale clusters in National Labs like Lawrence Livermore and Argonne National Labs and so on use this style of programming, using MPI as the message passing fabric. The only downside to the message passing style of programming is that it is difficult to program using this style. If you are a programmer who writes sequential programs, the transition path to writing an explicitly parallel program is easier if there is the notion of shared memory because it is natural to think of shared data structures among different threads of an application and that is the reason making the transition from sequential programming to parallel programming using for instance the `pthread` library on an SMP is a fairly intuitive and easy pathway. On the other hand, if the programmer has to think in terms of coordinating the activities on different processors by explicitly sending and receiving messages from the peers that is calling for a fairly radical change of thinking in terms of how to structure a program.

Distributed Shared Memory. That is the motivation for coming up this abstraction of distributed shared memory in a cluster. The idea is that we want to give the illusion to the application programmer writing an explicitly parallel program that all of the memory in the entire cluster is shared. They are not physically shared, but the DSM library is going to give the illusion to the threads running on each of these processors that all of the memory is shared. Thus, they will have an easier transition path for instance from a sequential program or going from a program that they write on an SMP to a program that runs on a cluster because they do not have to think in terms of message passing but they think in terms of shared memory. Also since we have provided shared memory semantics in DSM, there is no need for marshaling and unmarshaling arguments that are being passed from one processor to another and so on all of that is being handled by the fact there is "shared memory". So when you make a procedure call and that procedure call is touching some portion of memory that happens to be on remote memory. That memory is going to magically become available to the thread that is making that procedure call. In other words, the DSM abstraction gives the same level of comfort to a programmer who is used to programming a truly shared memory machine when they move to a cluster because they can use the same set of primitives like locks and barriers for synchronization and the `pthread` style of creating threads that will run on different nodes of the cluster and

Figure 4: History of hardware shared memory systems.


Figure 5: History of software DSM systems.

Figure 6: History of structured DSM systems.

that's the advantage of the DSM style of writing an explicitly parallel program.

3   History of Distributed Shared Memory Systems

Before going into details of DSM, let us first learn some history of the DSM systems. Figure 4, Figure 5, and Figure 6 give the history of hardware, software and structured distributed shared memory in the past 20 years, respectively. Note that, these figures only gave the approximate time of when these systems were invented. If you are interested in more details of the history of DSM systems, the survey in [26] may be a good starting point.

Figure 7: Shared memory synchronization.

4   Implementation of DSM

4.1   Shared Memory Programming: Revisit

In previous lectures, we have already learned the shared memory synchronization. **Lock** is a primitive and particularly the mutual exclusion lock is a primitive that is used ubiquitously in writing shared memory parallel programs to protect data structure so that one thread can exclusively modify the data and release the lock so that another thread can inspect the data later on. And similarly **barrier** synchronization is another synchronization primitive that is very popular in scientific programs.

In particularly, if you are writing a shared memory program, there are two types of memory accesses. One type of memory access is the normal reads and writes to shared data that is being manipulated by a particular thread. The second kind of memory access is for synchronization variables that are used in implementing locks and barriers by the operating system itself. It may be the operating system or it could be a user level threads library that is providing these mutual exclusion locks, or barriers primitives, but in implementing those synchronization primitives, those algorithms are going to use reads and writes to shared memory.

4.2   Memory Consistency

Recall that, in previous lectures, we have discussed the relationship between the memory consistency model and cache coherence, in the context of shared memory systems. Memory consistency model is a contract between the application programmer and the system. It answers the **when** question, that is, when a shared memory location is modified by one processor, when (*i.e.,* how soon) that change is going to be made visible to other processors that have the same memory location in their respective caches. Cache coherence, on the other hand, is answering the **how** question, that is, how is the system[1], implementing the contract of the memory consistency model ? In other words, the guarantee that is made by the memory consistency model,

---

[1]Note that, by system, we actually mean the system software plus the hardware working together.
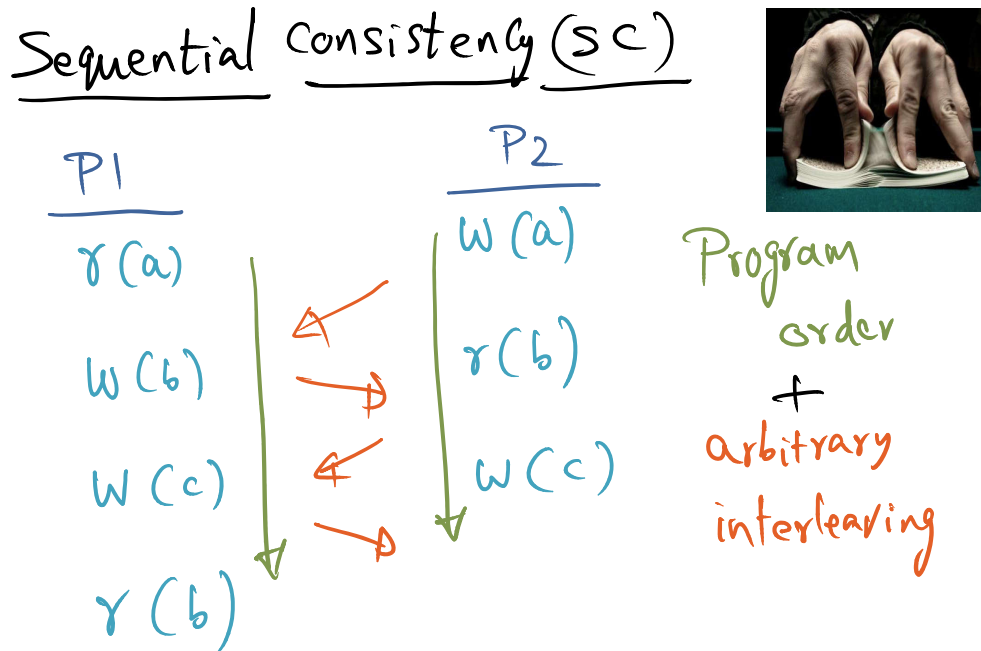
Figure 8: Sequential consistency.

to the application programmer has to be fulfilled by the cache coherence mechanism.

Now, coming back to writing a parallel program, when accesses are made to the shared memory, the underlying coherence mechanism has to ensure that all the processors see the changes that are being made to shared memory, commensurate with the memory consistency model.

Sequential Consistency (SC). One particular memory consistency model that we have discussed in previous lectures is the sequential consistency. The idea behind sequential consistency is very simple. That is, every process is making some memory accesses. For example, as shown in Figure 8, Process $P1$ have 4 memory accesses. From the perspective of the programmer, the expectation is that, these memory accesses are happening in the textual order (*i.e.,* read(a) $\rightarrow$ write(b) $\rightarrow$ write(c) $\rightarrow$ read(b)). Similarly, the expectation of the memory accesses of Process $P2$ is that write(a)$\rightarrow$ read(b) $\rightarrow$ write(c). The real question is what happens to the accesses that are happening on one processor with respect to the accesses on another processor if they are accessing exactly the same memory location ? For instance, Process $P1$ is reading location a, and Process $P2$ is writing location a. What is the order between this read by $P1$ and this write by $P2$? That is where sequential consistency model says that the interleaving of memory accesses between multiple processors (Figure 8 is showing two, but you can have arbitrary number of those processors) making accesses to shared memory all in parallel. When that happens you want to observe the textual program order for the accesses in an individual processor but the interleaving of the memory accesses coming from the different processors is arbitrary. In other words, the sequential memory consistency model builds on the atomicity for individual read/write operations and says that, individual read/write operations are atomic on any given processor, and the program order has to be preserved. And, in order to think about interleaving of the memory accesses that are happening on different processors. That can be arbitrary and should be consistent with the thinking of the programmer.

Back to our parallel program. As shown in Figure 9, a parallel program is making read/write accesses to shared memory, some of them offer data, some of them offer synchronization. Now, as far as the sequentially consistent memory model is concerned, it does not distinguish between accesses coming from processors as data accesses, or synchronization accesses. So there will be coherence action on every read/write access that the model sees. For example, if $P1$ (in Figure 9) writes to a memory location, then the sequentially
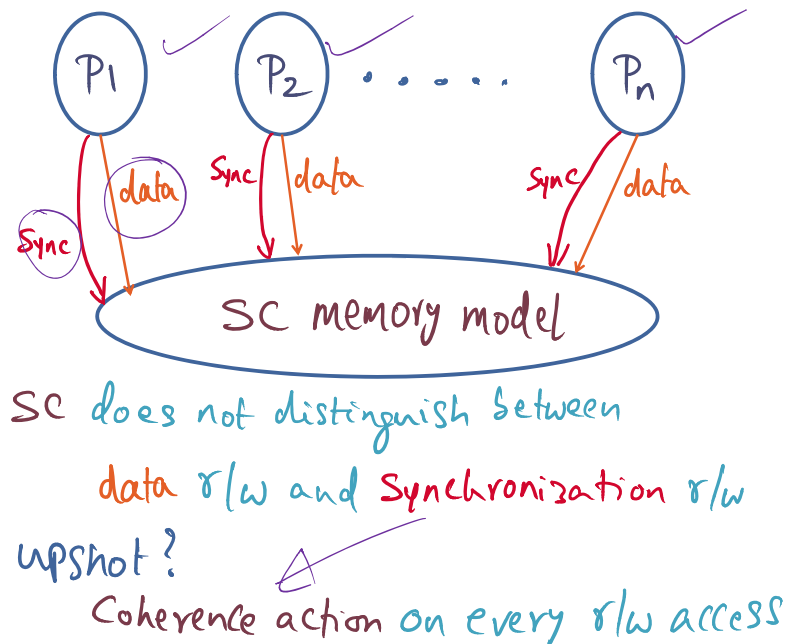
Figure 9: Illustration of sequential consistency in a parallel program.

consistent memory model has to ensure that this write is inserted into the global order somewhere. However, to insert into the global order somewhere, it has to perform the coherence action with respect to all the other processors. That is the upshot of not distinguishing between normal data accesses and synchronization accesses that is inherent in the SC memory model.

Now, let us see a typical parallel program shown in Figure 10. In the program, we decided that the accesses to shared variables `a` and `b` should be governed by a lock `L`. So if any process wants to read or write variables `a` and `b`, it will get a lock and then mess with variables that are governed by this lock. Once the process is done with whatever it wants to do with these shared variables, it will unlock indicating that it is done. And the codes between `Lock(L)` and `Unlock(L)` are the critical section. Within the critical section, processes are allowed to do whatever they want on these data structures that are governed by this particular lock, because that is an association we as the programmer have made in writing the parallel program. So if another process say $P2$, gets the same lock. It is going to get the lock only after $P1$ has released it. And consequently, if you look at the structure of the critical section of $P2$, it gets a lock and it is messing with the same data structures that $P1$ was messing with in its critical section. But, by design, we know that either $P1$ or $P2$ can be messing with the data structures at any point of time. And that is a guarantee that we know comes from the fact that we designed the parallel program. And the lock is associated with these data structures. So in other words, $P2$ is not going to access any of the data that is inside the critical section until $P1$ releases the lock. We know this because we design the program, but the SC memory model does not know about the association between these data structures and the lock. And in particular, it does not even know that memory access emanating from the processor due to the lock primitive is a "different animal" compared to the memory accesses coming from the processor as a results of accessing normal data structures. So the cache coherence mechanism that is provided by the system for implementing the memory consistency model is going to be doing more work that it needs to do because it is going to take actions on every one of these accesses, even though the coherence actions are not warranted for these data in the critical section of $P1$ until $P1$ releases the lock. So what that means is that there is going to be more overhead for maintaining the coherence commensurate with the SC memory model. That means it is going to lead to a pooler scalability of the shared memory system. So in this particular example, since $P2$ is not going to access any of these data structures in its critical section until $P1$ has released the lock. There is no need for coherence action
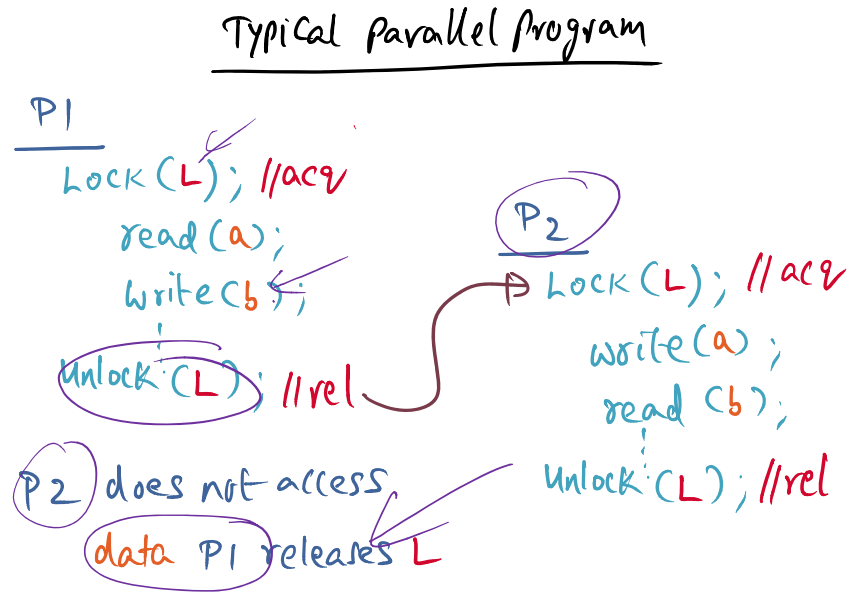
Figure 10: A typical parallel program.

for `a` and `b` until the lock is actually released.

Release Consistency (RC). This is the motivation for another memory consistency model, which is called release consistency. Generally, the parallel program consists of several different parallel threads (*e.g.,* $P1$ and $P2$ in Figure 11), and if any process (or thread) wants to mess with some shared data structures, it is going to acquire a lock (*e.g., a*1 in Figure 11) and in the mind of the programmer there is an association between this lock and the data structures governed by it. So long as they hold the lock they can modify the data structure and then release the lock (*e.g., r*1 in Figure 11). So every critical section you can think of as composed of an acquire followed by data accesses governed by the lock and then a release. If the same lock is used by some other process (*e.g., P*2) and if the critical section of $P1$ preceded that of $P2$. In other words, $r1$ of $P1$ happens before $a2$ of $P2$. If this acquire operation for the same lock happened after the release of that lock. All that we have to ensure is that all the coherence actions prior to this release of the lock by $P1$ has to be complete before we allow $P2$ to acquire this lock. That is the idea of release consistency. So we take the synchronization operations that are provided by the system whether it is hardware or software and we label them as either an acquire operation or a release operation. So it is very straightforward when you think about mutual exclusion lock where the lock primitive is the acquire operation, while the unlock primitive is the release operation. Other synchronization operations can also be mapped to acquire and release, for example, the barrier synchronization. The arrival of a barrier is equivalent to an acquire and leaving the barrier is equivalent to a release.

Therefore, if $P1$ does a shared memory access within its critical section, and that shared memory access would normally result in some coherence actions on the interconnect reaching to the other processes and so on. If it is a SC memory model we will block process $P1$ until that particular memory access is complete with respect to all the processes in the shared memory machine. But if we use the release consistency model, we do not have to block $P1$ in order for coherence actions to be complete to let the process continue on with its computation. We only have to block a process at a release point to make sure that any coherence actions that may have been initiated up until the release point are all complete before we perform this release operation. So the RC model allows exploitation of computation on $P1$ with communication that may be happening through the coherence mechanism for completing the coherence actions corresponding to the
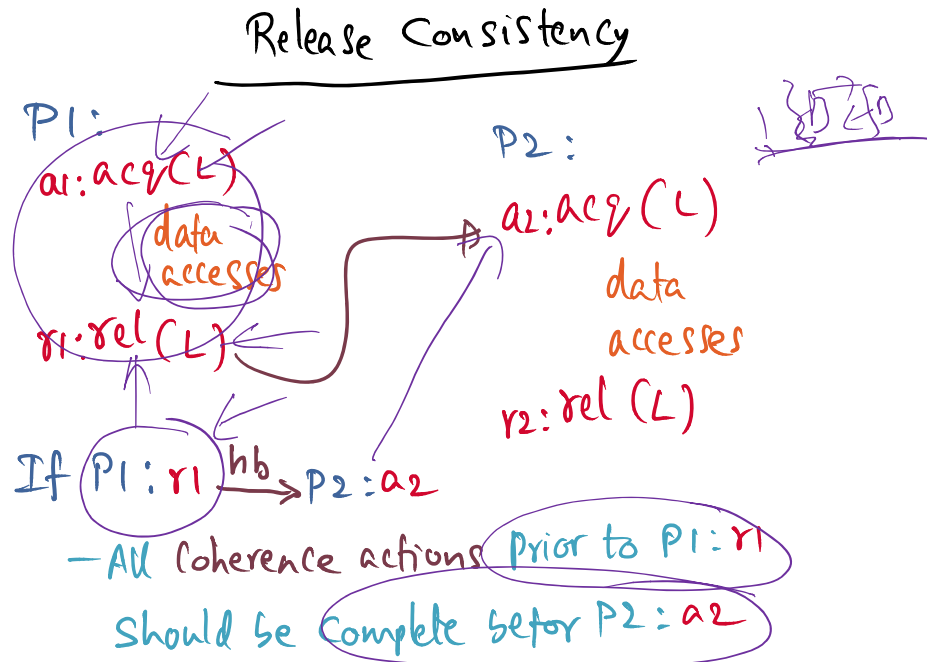
Figure 11: Release consistency.

memory accesses that are making inside the critical section.

Lazy RC. Now, let us see a lazy version of the RC memory model. It is called LRC. Figure 11 gives you the structure of a parallel program: a thread acquires a lock, does the data accesses, releases the lock; another thread may acquire the same lock. Recall that if the critical section of $P1$ precedes that of $P2$, then the RC memory model requires that, at the point of release (*i.e.*, $r1$), you ensure that all the modifications that have been made on processor $P1$ are all communicated to other processors through coherence actions. Then, $P1$ releases the lock. That is called **eager release consistency**, meaning that at the point of release you are ensuring that the entire system is cache coherent at the point of release. Assume that $P2$'s acquire happens much later than $P1$'s release of the same lock. Clearly, there is an opportunity for procrastination. We have already seen that procrastination often helps in system design (*e.g.*, mutual exclusion locks and process scheduling) in previous lectures. So lazy RC is another instance where **procrastination may actually help in optimizing the system performance**. The idea is that, rather than performing all the coherence actions at the point of release. But wait till the acquire actually happens. At the point of acquire, take all the coherence actions before allowing this acquire to succeed. So the key point is that you are deferring the point at which you ensure that all the coherence actions are complete to the point of acquisition as opposed to the point of release.

Figure 12 gives the comparison between eager RC and lazy RC. As you can see, in the eager version of RC memory model, $P1$ has to communicate to all other processors to ensure that the changes on variable x are visible to them at the release point, and so does $P2$. Now move over to the lazy version. Note that, in the lazy version, we did nothing while $P1$ releases its lock, later on, the next processor that happens to acquire the same lock. The RC memory model has to make sure that all coherence actions associated with that particular lock have been complete. In this example, the previous lock holder, *i.e.*, $P1$, has made changes to variable x, so the RC memory model will pull it from $P1$, and then $P2$ can execute its critical section. Then, before $P3$ executes its critical section, it is going to pull it from $P2$, and complete what it needs to do. The most important thing that we need to pay attention to is that there is no broadcast anymore. There is only point-to-point communication that is happening between processors that are passing the lock from one to the other. Note that, lazy RC memory model is also called **pull model**, since we are pulling the coherence

Figure 12: Eager RC vs. lazy RC.

Table 1: Eager RC vs. lazy RC.

Eager RC		Lazy RC	
Pros	**Cons**	**Pros**	**Cons**
Lower latency at acquisition	More messages	Less messages	More latency at acquisition

actions at the point of acquisition. Whereas the eager version is called **push model**. Table 1 summarizes the pros and cons of the eager RC and lazy RC, respectively.

4.3   Software DSM

So far, we have seen three types of memory consistency model: sequential consistency, eager release consistency, and lazy release consistency. Strictly specking, the latter two are just two different versions of the release consistency memory model. Now it is the time to take the transition to talking about **software distributed shared memory** and we will see how these memory models are coming to play in building software distributed memory.

Recall that, we are dealing with a computational cluster. In the cluster, each node has its own private physical memory, but there is no physically shared memory. Therefore, the system software has to implement the consistency model to the programmer. In a tightly coupled multiprocessor system, coherence is maintained at individual memory access level by the hardware. Unfortunately, that fine-grain coherence maintenance at individual memory access level will lead to too much overhead in a cluster. Why ? Because on every load or store instruction that is happening on any one of these processors, the system software has to butt in, and implement the coherence actions in software through the entire cluster. Clearly, that is infeasible. So what do we do to implement software DSM?

The first thought is to implement this sharing and coherence maintenance at a much coarse level, for example,

at the level of pages. Actually, in an individual processor, or shared memory multiprocessor system, the coherence maintenance is not simply a single word that a processor is doing a load or a store on. Because in order to exploit spatial locality, the block size used in caches in processors tend to be bigger than the granularity of memory access that is possible from individual instructions in the processor. Here we are going to do the coherence maintenance in software, and we keep the granularity of coherence maintenance to be an entire page. And we are going to maintain the coherence of the distributed shared memory in software by cooperating with the operating system that is running on every node. So what we are going to do is we are providing a global virtual memory abstraction to the application programmer running on the cluster. So the application programmer can view the entire cluster as a globally shared virtual memory system. Under the cover what the DSM software is doing is to partition the global address space into chunks that are managed individually on the nodes of the different processors of the cluster. From the application point of view, what this global virtual memory abstraction is giving is address equivalence and that is if I access a memory location, say X in the program that means exactly the same thing whether I access the memory location X from processor 1 or processor 2 and so on so forth. That is the idea in providing a global virtual memory abstraction. And the way the DSM software is going to handle maintenance of coherence is by having distributed ownership for the different virtual pages that constitute this global virtual address space. So you can think of the global virtual memory space as constituted by several pages. And we are going to say some number of pages are owned by processor 1, some number of pages are owned by processor 2, and so on so forth. That is we split the ownership responsibility into individual processors. What that means is that the owner of a particular page is also responsible for keeping complete coherence information for that particular page and taking the coherence actions commensurate with that page. And the local physical memory available in each processor is being used for hosting portions of the global virtual memory space in the individual processor commensurate with the access pattern that is being displayed by the application on the different processors. For instance, as shown in Figure 13, if processor 1 accesses certain portion of the global virtual memory space, then this portion is mapped into the local physical memory of this processor. So that the thread that is running one this processor can access this portion of the global address space. And it might be that the same page is being shared with some other processor. In that case, a copy of this page is existing in both of the two processors. Now it is up to the processor that is responsible for the ownership of this particular page to worry about the consistency of this page that is now resident in multiple locations. For instance, if node 1 (*i.e.,* $P1$) is the owner of this page, then node 1 will have the metadata that indicates that this particular page currently shared by $P1$ and PX. That is the directory that is associated with the portion of the global virtual memory space that is being owned and managed by $P1$. So statically we make an association between a portion of the address space and the owner for that portion of the address space in terms of coherence maintenance for that portion of the global virtual memory space.

Single-Writer Protocol for Software DSM. Now it is the time to dive into the software DSM shown in Figure 13. As you can see there are several layers in the software DSM. The one between the blue line and the green line is the abstraction layer seen by the application, which is giving this illusion of a global virtual memory. The next layer (*i.e.,* the next one below the abstraction layer) is the DSM software implementation layer, which implements this global virtual memory abstraction. In particular, this DSM software layer, which exists on every processor, knows that the point of access to a page by a processor, who exactly to contact, as the owner of the page, to get the current copy of the page. For instance, suppose there was a page fault on processor 1. For a particular portion of the global address space. That portion of the global address space is currently not resident in $P1$'s local physical memory. So there is a page fault, and there is cooperation, as mentioned earlier, between the operating system and the DSM software. So when the page fault happens, the page fault is going to be communicated by the operating system to the DSM software. What the DSM software is going to do is, it knows the owner f the page, and so it is going to contact the owner of the page. And ask the owner to get the current copy of the page. Note that, the owner itself might not have the current copy of the page, but it knows which node has, and if that is the case, the owner is going to fetch the current copy of the page and send over to the node that is requesting it. Once the page is brought in the physical memory, the DSM software contacts the virtual memory manager notifying that it has completed and asking the virtual memory manager to update the page table of that processor so that it can resume execution. Then, the VM manager gets into action and updates the page table for the thread to indicate that the faulty virtual page is now mapped to a physical page. And then the processor or the thread
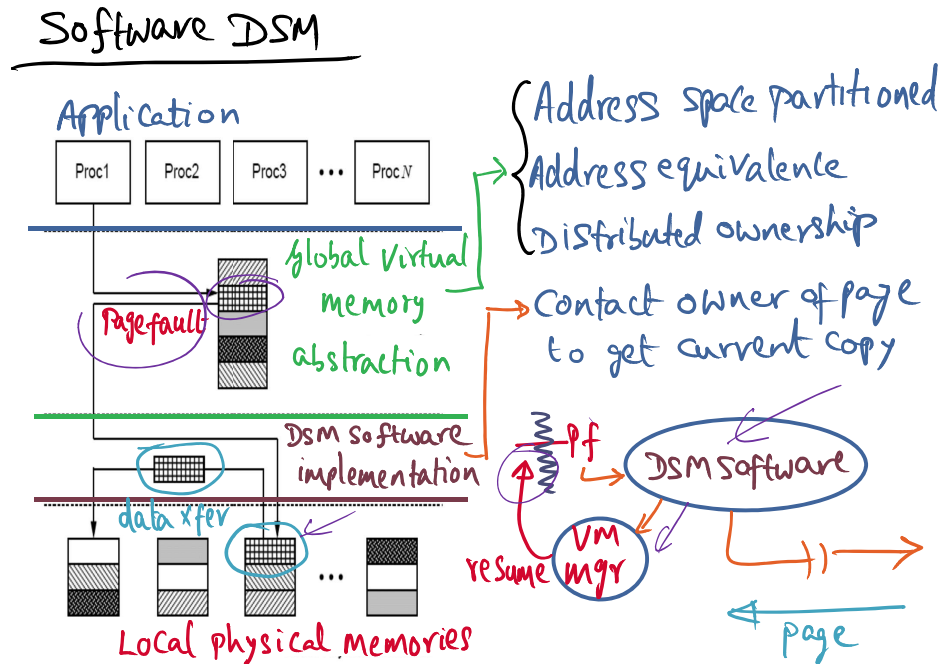
Figure 13: Software DSM.

can resume its execution. So that is the way coherence is going to be maintained by the DSM software and the cooperation between DSM software and the VM manager. And the coherence maintenance is happening at the level of individual pages.

Early examples of systems that built software DSM include IVY from Yale [12], CLOUDS from Georgia Tech. [13], Mirage from UPenn [14], and Munin from Rice [16]. All of these are distributed shared systems and they all used coherence maintenance at the granularity of an individual page, and they used a protocol which is often referred to as a single-writer protocol. That is, as mentioned that the directory associated with the portion of the virtual memory space managed by each one of the nodes. And the directory has information as to who all are sharing a page at any point of time. Multiple readers can share a page at any point of time, but a single writer is only allowed to have the page at any point of time. So if there is a writer for a particular page, say that page is now currently in the memory of two different processors, if the writer wants to write to this page, then it has to inform through the DSM software abstraction, inform the owner of this page. Then, the owner is going to invalidate all copies of this page that exists in the entire system. So that this writer has exclusive access to that page so that it can make modifications to it. That is the single writer multiple reader protocol.

Unfortunately, there is a "problem" with the single writer protocol. That is there is potential for what is called false sharing, which we have talked about already in the context of shared memory multiprocessor systems. Basically the concept of false sharing is that data appears to be shared even though programmatically they are not. Considering this page-based coherence maintenance. In the page-based coherence, the coherence maintenance that is done by the software, i.e., DSM software, is of the granularity of a single page. A page may be $4K$ or $8K$ bytes, depending on the architecture. Clearly, within a page, lots of different data structures can actually fit. So if the coherence maintenance is being done at the level of an individual page, then we are invalidating copies of the page in several nodes to allow one writer to make modifications to some portions of the page. And that can be very severe in a page based system due the coarse granularity of the coherence information. So for example, this one page may contain ten different data structures, each of which is governed by a distinct lock as far as the application programmer is concerned. But even if I get a lock, which is for a particular data structure that happens to be in this page. And some other processor

has a lock for a different data structure which is also on this page. When I get the lock that is going to manipulate a particular data structure in this page, and if I want to make modifications to it, I am going to go and invalidate all the other copies. When the other processor wants to make a change, it is going to come and invalidate my copy of this page. So the page can be ping-pong between multiple processors. Even though they are modifying different portions of the same page. Still the coherence granularity being a page, will results in this page shuttling back and forth between these two processors, even though the application program is perfectly well behaved in terms of using locks to govern different data structures. Unfortunately, all of the data structures happened to fit within the same page, resulting in this false sharing. Therefore, page-level granularity and single writer multiple reader protocol do not live happily together. They will lead to false sharing and ping-pong of the pages due to the false sharing among the threads of the application across the entire network.

In the next lecture, we will cover a solution to mitigate the false sharing issue of the single writer multiple reader protocol, and some implementation details.

References

[1] "Parallel Virtual Machine." [Online]. Available: http://www.csm.ornl.gov/pvm/

[2] K. Watkins, "CLF: A logical framework for concurrent systems," 2003.

[3] J. S. Alherbish, *The BBN butterfly family: An overview.* Berlin, Heidelberg: Springer Berlin Heidelberg, 1994, pp. 489–490. [Online]. Available: https://doi.org/10.1007/3-540-57981-8_166

[4] G. Graunke and S. Thakkar, "Synchronization algorithms for shared-memory multiprocessors," *Computer*, vol. 23, no. 6, pp. 60–69, June 1990.

[5] U. Ramachandran, G. Shah, S. Ravikumar, and J. Muthukumarasamy, "Scalability study of the ksr-1," in *Parallel Processing, 1993. ICPP 1993. International Conference on*, vol. 1. IEEE, 1993, pp. 237–240.

[6] A. Agarwal, D. Chaiken, K. Johnson, D. Kranz, J. Kubiatowicz, K. Kurihara, B. H. Lim, G. Maa, D. Nussbaum, M. Parkin, and D. Yeung, "The mit alewife machine: A large-scale distributed-memory multiprocessor," Cambridge, MA, USA, Tech. Rep., 1991.

[7] D. Lenoski, J. Laudon, K. Gharachorloo, W.-D. Weber, A. Gupta, J. Hennessy, M. Horowitz, and M. S. Lam, "The stanford dash multiprocessor," *Computer*, vol. 25, no. 3, pp. 63–79, Mar. 1992. [Online]. Available: http://dx.doi.org/10.1109/2.121510

[8] "SGI 2000." [Online]. Available: https://en.wikipedia.org/wiki/SGI_Origin_2000

[9] "SGI Altix." [Online]. Available: https://en.wikipedia.org/wiki/Altix

[10] "IBM Blue Gene." [Online]. Available: https://en.wikipedia.org/wiki/Blue_Gene

[11] D. R. Helman, "Sorting on clusters of smps," in *Proceedings of the First Merged International Parallel Processing Symposium and Symposium on Parallel and Distributed Processing*, Mar 1998, pp. 561–567.

[12] K. Li and P. Hudak, "Memory coherence in shared virtual memory systems," *ACM Trans. Comput. Syst.*, vol. 7, no. 4, pp. 321–359, Nov. 1989. [Online]. Available: http://doi.acm.org/10.1145/75104.75105

[13] A. Mohindra, "Issues in the design of distributed shared memory systems," Georgia Institute of Technology, Tech. Rep., 1993.

[14] B. Fleisch and G. Popek, "Mirage: A coherent distributed shared memory design," *SIGOPS Oper. Syst. Rev.*, vol. 23, no. 5, pp. 211–223, Nov. 1989. [Online]. Available: http://doi.acm.org/10.1145/74851.74871

[15] R. G. Minnich and D. V. Pryor, "Mether: supporting the shared memory model on computing clusters," in *Digest of Papers. Compcon Spring*, Feb 1993, pp. 558–567.

[16] J. K. Bennett, J. B. Carter, and W. Zwaenepoel, "Munin: Distributed shared memory based on type-specific memory coherence," *SIGPLAN Not.*, vol. 25, no. 3, pp. 168–176, Feb. 1990. [Online]. Available: http://doi.acm.org/10.1145/99164.99182

[17] C. Amza, A. L. Cox, S. Dwarkadas, P. Keleher, H. Lu, R. Rajamony, W. Yu, and W. Zwaenepoel, "Treadmarks: shared memory computing on networks of workstations," *Computer*, vol. 29, no. 2, pp. 18–28, Feb 1996.

[18] I. Schoinas, B. Falsafi, A. R. Lebeck, S. K. Reinhardt, J. R. Larus, and D. A. Wood, "Fine-grain access control for distributed shared memory," *SIGPLAN Not.*, vol. 29, no. 11, pp. 297–306, Nov. 1994. [Online]. Available: http://doi.acm.org/10.1145/195470.195575

[19] D. J. Scales and K. Gharachorloo, "Design and performance of the shasta distributed shared memory protocol," in *Proceedings of the 11th International Conference on Supercomputing*, ser. ICS '97. New York, NY, USA: ACM, 1997, pp. 245–252. [Online]. Available: http://doi.acm.org/10.1145/263580.263643

[20] S. Dwarkadas, N. Hardavellas, L. Kontothanassis, R. Nikhil, and R. Stets, "Cashmere-vlm: Remote memory paging for software distributed shared memory," in *Proceedings 13th International Parallel Processing Symposium and 10th Symposium on Parallel and Distributed Processing. IPPS/SPDP 1999*, Apr 1999, pp. 153–159.

[21] "The Beehive Cluster System." [Online]. Available: https://www.cc.gatech.edu/~rama/Beehive/

[22] S. Ahuja, N. Carriero, and D. Gelernter, "Linda and friends," *Computer*, vol. 19, no. 8, pp. 26–34, Aug 1986.

[23] H. E. Bal, M. F. Kaashoek, and A. S. Tanenbaum, "Orca: A language for parallel programming of distributed systems," *IEEE transactions on software engineering*, vol. 18, no. 3, pp. 190–205, 1992.

[24] R. S. Nikhil, U. Ramachandran, J. M. Rehg, R. H. Halstead, C. F. Joerg, and L. Kontothanassis, "Stampede a programming system for emerging scalable interactive multimedia applications," in *International Workshop on Languages and Compilers for Parallel Computing*. Springer, 1998, pp. 83–99.

[25] D. Hilley and U. Ramachandran, "Stampede rt: Programming abstractions for live streaming applications," in *27th International Conference on Distributed Computing Systems (ICDCS '07)*, June 2007, pp. 65–65.

[26] B. Nitzberg and V. Lo, "Distributed shared memory: a survey of issues and algorithms," *Computer*, vol. 24, no. 8, pp. 52–60, Aug 1991.